



Photo credit: Pixabay

Multi-Class Text Classification Model Comparison and Selection

Natural Language Processing, word2vec, Support Vector Machine, bag-of-words, deep learning

Susan Li [Follow](#)

Sep 25, 2018 · 7 min read

When working on a supervised machine learning problem with a given data set, we try different algorithms and techniques to search for models to produce general hypotheses, which then make the most accurate predictions possible about future instances. The same principles apply to text (or document) classification where there are many models can be used to train a text classifier. The answer to the question “What machine learning model should I use?” is always “It depends.” Even the most experienced data scientists can’t tell which algorithm will perform best before experimenting them.

This is what we are going to do today: use everything that we have presented about text classification in the previous articles (and more) and comparing between the text classification models we trained in order to choose the most accurate one for our problem.

The Data

We are using a relatively large data set of Stack Overflow questions and tags. The data is available in Google BigQuery, it is also publicly

available at this Cloud Storage URL:

<https://storage.googleapis.com/tensorflow-workshop-examples/stackoverflow-data.csv>.

Exploring the Data

```

1  import logging
2  import pandas as pd
3  import numpy as np
4  from numpy import random
5  import gensim
6  import nltk
7  from sklearn.model_selection import train_test_split
8  from sklearn.feature_extraction.text import CountVector
9  from sklearn.metrics import accuracy_score, confusion_
10 import matplotlib.pyplot as plt
11 from nltk.corpus import stopwords
12 import re
13 from bs4 import BeautifulSoup

```

explore

	post	tags
0	what is causing this behavior in our c# datet...	c#
1	have dynamic html load as if it was in an ifra...	asp.net
2	how to convert a float value in to min:sec i ...	objective-c
3	.net framework 4 redistributable just wonderi...	.net
4	trying to calculate and print the mean and its...	python
5	how to give alias name for my website i have ...	asp.net
6	window.open() returns null in angularjs it wo...	angularjs
7	identifying server timeout quickly in iphone ...	iphone
8	unknown method key error in rails 2.3.8 unit ...	ruby-on-rails
9	from the include how to show and hide the con...	angularjs

Figure 1

10276752

We have over 10 million words in the data.

```

my_tags = ['java','html','asp.net','c#','ruby-on-
rails','jquery','mysql','php','ios','javascript','python','c
','css','android','iphone','sql','objective-

```

```
c','c++','angularjs','.net']
plt.figure(figsize=(10,4))
df.tags.value_counts().plot(kind='bar');
```




Figure 2

The classes are very well balanced.

We want to have a look a few post and tag pairs.

```
def print_plot(index):
    example = df[df.index == index][['post',
    'tags']].values[0]
    if len(example) > 0:
        print(example[0])
        print('Tag:', example[1])
```

```
print_plot(10)
```

```
when we need interface c# <blockquote> <strong>possible duplicate:</strong><br> <a href= https://stackoverflow.com/question/240152/why-would-i-want-to-use-interfaces >why would i want to use interfaces </a> <a href= https://stackoverflow.com/questions/9451868/why-i-need-interface >why i need interface </a> </blockquote> i want to know where and when to use it
for example <pre><code>interface idemo { // function prototype public void show(); } // first class using the interface c
lass myclass1 : idemo { public void show() { // function body comes here response.write( i m in myclass ); } } // seco
nd class using the interface class myclass2 : idemo { public void show() { // function body comes here response.write( i
m in myclass2 ); response.write( so what ); } </code></pre> these two classes has the same function name with different
body. this can be even achieved without interface. then why we need an interface where and when to use it
Tag: c#
```

Figure 3

```
print_plot(30)
```

how to chain expressions inside ngclass when using the {...}[] form how can i add another expression to an `<code>ng-class</code>`

directive that uses this form: `<pre><code>ng-class= {true: loading false: loading-done }[data.loader===null] </code></pre>`

i d like to add something like this to the list: `<pre><code>{highlight:isspecial} </code></pre>` is it possible without expanding the first expression thanks.

Tag: angularjs

Figure 4

As you can see, the texts need to be cleaned up.

Text Pre-processing

The text cleaning techniques we have seen so far work very well in practice. Depending on the kind of texts you may encounter, it may be relevant to include more complex text cleaning steps. But keep in mind that the more steps we add, the longer the text cleaning will take.

For this particular data set, our text cleaning step includes HTML decoding, remove stop words, change text to lower case, remove punctuation, remove bad characters, and so on.

```
1  REPLACE_BY_SPACE_RE = re.compile('[/(){}\\[\\]\\|@,;]')
2  BAD_SYMBOLS_RE = re.compile('[^0-9a-z #+_]')
3  STOPWORDS = set(stopwords.words('english'))
4
5  def clean_text(text):
6      """
7          text: a string
8
9          return: modified initial string
10     """
11     text = BeautifulSoup(text, "lxml").text # HTML dec
12     text = text.lower() # lowercase text
13     text = REPLACE_BY_SPACE_RE.sub(' ', text) # replac
```

clean_text

Now we can have a look a cleaned post:

need interface c# possible duplicate would want use interfaces need interface want know use example interface idemo function pr
 otype public void show first class using interface class myclass1 idemo public void show function body comes responsewrite my
 class second class using interface class myclass2 idemo public void show function body comes responsewrite myclass2 responsewri
 te two classes function name different body even achieved without interface need interface use
 Tag: c#

Figure 5

Way better!

```
df['post'].apply(lambda x: len(x.split(' '))).sum()
```

3421180

After text cleaning and removing stop words, we have only over 3 million words to work with!

After splitting the data set, the next steps includes feature engineering. We will convert our text documents to a matrix of token counts (CountVectorizer), then transform a count matrix to a normalized tf-idf representation (tf-idf transformer). After that, we train several classifiers from [Scikit-Learn library](#).

```
X = df.post
y = df.tags
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.3, random_state = 42)
```

Naive Bayes Classifier for Multinomial Models

After we have our features, we can train a classifier to try to predict the tag of a post. We will start with a [Naive Bayes](#) classifier, which provides a nice baseline for this task. `scikit-learn` includes several variants of this classifier; the one most suitable for text is the multinomial variant.

To make the vectorizer => transformer => classifier easier to work with, we will use `Pipeline` class in Scikit-Learn that behaves like a compound classifier.

```

1  from sklearn.naive_bayes import MultinomialNB
2  from sklearn.pipeline import Pipeline
3  from sklearn.feature_extraction.text import TfidfTransformer
4
5  nb = Pipeline([('vect', CountVectorizer()),
6                 ('tfidf', TfidfTransformer()),
7                 ('clf', MultinomialNB()),
8                 ])
9  nb.fit(X_train, y_train)
10
11  %%time

```

nb

```

accuracy 0.7395
          precision    recall  f1-score   support

   java         0.63      0.65      0.64       613
   html         0.94      0.86      0.90       620
  asp.net         0.87      0.92      0.90       587
    c#          0.70      0.77      0.73       586
ruby-on-rails     0.73      0.87      0.79       599
   jquery        0.72      0.51      0.60       589
   mysql         0.77      0.74      0.75       594
   php          0.69      0.89      0.78       610
   ios          0.63      0.59      0.61       617
 javascript       0.57      0.65      0.61       587
   python        0.70      0.50      0.59       611
    c           0.79      0.79      0.79       594
    css          0.84      0.59      0.69       619
  android        0.66      0.84      0.74       574
  iphone         0.64      0.83      0.72       584
    sql         0.66      0.64      0.65       578
objective-c       0.79      0.77      0.78       591
    c++          0.89      0.83      0.86       608
 angularjs       0.94      0.89      0.91       638
    .net         0.74      0.66      0.70       601

 avg / total         0.75      0.74      0.74    12000

Wall time: 955 ms

```

Figure 6

We achieved 74% accuracy.

Linear Support Vector Machine

Linear Support Vector Machine is widely regarded as one of the best text classification algorithms.

```

1  from sklearn.linear_model import SGDClassifier
2
3  sgd = Pipeline([('vect', CountVectorizer()),
4                  ('tfidf', TfidfTransformer()),
5                  ('clf', SGDClassifier(loss='hinge', pe
6                  )])
7  sgd.fit(X_train, y_train)
8
9  %%time
10

```

svm

```

accuracy 0.7891666666666667
          precision    recall  f1-score   support

   java      0.74      0.68      0.71      613
   html      0.85      0.93      0.89      620
  asp.net      0.87      0.95      0.91      587
    c#       0.81      0.80      0.80      586
ruby-on-rails 0.74      0.88      0.80      599
   jquery 0.77      0.41      0.53      589
   mysql    0.82      0.68      0.74      594
    php     0.70      0.95      0.81      610
    ios     0.82      0.56      0.66      617
 javascript 0.72      0.59      0.65      587
   python 0.71      0.65      0.68      611
    c      0.81      0.87      0.84      594
    css     0.77      0.79      0.78      619
  android 0.83      0.86      0.85      574
  iphone    0.81      0.80      0.81      584
    sql     0.71      0.68      0.69      578
objective-c 0.81      0.90      0.85      591
    c++     0.84      0.96      0.89      608
 angularjs 0.87      0.95      0.91      638
    .net    0.77      0.89      0.83      601

 avg / total      0.79      0.79      0.78     12000

Wall time: 1.26 s

```

Figure 7

We achieve a higher accuracy score of 79% which is 5% improvement over Naive Bayes.

Logistic Regression

Logistic regression is a simple and easy to understand classification algorithm, and Logistic regression can be easily generalized to multiple classes.

```

1  from sklearn.linear_model import LogisticRegression
2
3  logreg = Pipeline([('vect', CountVectorizer()),
4                     ('tfidf', TfidfTransformer()),
5                     ('clf', LogisticRegression(n_jobs=1, C=1e5)),
6                     ])
7  logreg.fit(X_train, y_train)
8
9  %%time
10

```

logreg

```

accuracy 0.783
precision recall f1-score support
java      0.70    0.62    0.66    613
html      0.91    0.91    0.91    620
asp.net   0.97    0.94    0.95    587
c#        0.78    0.77    0.78    586
ruby-on-rails 0.77    0.81    0.79    599
jquery    0.59    0.58    0.58    589
mysql     0.77    0.76    0.76    594
php       0.82    0.86    0.84    610
ios       0.70    0.72    0.71    617
javascript 0.61    0.59    0.60    587
python    0.64    0.63    0.64    611
c         0.83    0.83    0.83    594
css       0.78    0.78    0.78    619
android   0.85    0.85    0.85    574
iphone    0.80    0.83    0.81    584
sql       0.65    0.65    0.65    578
objective-c 0.82    0.84    0.83    591
c++       0.91    0.91    0.91    608
angularjs 0.96    0.94    0.95    638
.net      0.78    0.83    0.80    601

avg / total    0.78    0.78    0.78    12000

Wall time: 981 ms

```

Figure 8

We achieve an accuracy score of 78% which is 4% higher than Naive Bayes and 1% lower than SVM.

As you can see, following some very basic steps and using a simple linear model, we were able to reach as high as an 79% accuracy on this multi-class text classification data set.

Using the same data set, we are going to try some advanced techniques such as word embedding and neural networks.

Now, let's try some complex features than just simply counting words.

Word2vec and Logistic Regression

Word2vec, like doc2vec, belongs to the text preprocessing phase. Specifically, to the part that transforms a text into a row of numbers.

Word2vec is a type of mapping that allows words with similar meaning to have similar vector representation.

The idea behind Word2vec is rather simple: we want to use the surrounding words to represent the target words with a Neural Network whose hidden layer encodes the word representation.

First we load a word2vec model. It has been pre-trained by Google on a 100 billion word Google News corpus.

```
from gensim.models import Word2Vec

wv =
gensim.models.KeyedVectors.load_word2vec_format("GoogleNews-
vectors-negative300.bin.gz", binary=True)
wv.init_sims(replace=True)
```

We may want to explore some vocabularies.

```
from itertools import islice
list(islice(wv.vocab, 13030, 13050))
```

```
['Memorial_Hospital',
 'Seniors',
 'memorandum',
 'elephant',
 'Trump',
 'Census',
 'pilgrims',
 'De',
 'Dogs',
 '###-####_ext',
 'chaotic',
 'forgive',
 'scholar',
 'Lottery',
 'decreasing',
 'Supervisor',
 'fundamentally',
 'Fitness',
 'abundance',
 'Hold']
```

Figure 9

BOW based approaches that includes averaging, summation, weighted addition. The common way is to average the two word vectors. Therefore, we will follow the most common way.

```

1  def word_averaging(wv, words):
2      all_words, mean = set(), []
3
4      for word in words:
5          if isinstance(word, np.ndarray):
6              mean.append(word)
7          elif word in wv.vocab:
8              mean.append(wv.syn0norm[wv.vocab[word].index])
9              all_words.add(wv.vocab[word].index)
10
11     if not mean:
12         logging.warning("cannot compute similarity with")
13         # FIXME: remove these examples in pre-processing
14         return np.zeros(wv.vector_size)

```

word_averaging

We will tokenize the text and apply the tokenization to “post” column, and apply word vector averaging to tokenized text.

```

1  def w2v_tokenize_text(text):
2      tokens = []
3      for sent in nltk.sent_tokenize(text, language='eng'):
4          for word in nltk.word_tokenize(sent, language='eng'):
5              if len(word) < 2:
6                  continue
7              tokens.append(word)
8      return tokens
9
10 train, test = train_test_split(df, test_size=0.3, random_state=42)
11

```

w2v_tokenize_text

Its time to see how logistic regression classifiers performs on these word-averaging document features.

```

from sklearn.linear_model import LogisticRegression
logreg = LogisticRegression(n_jobs=1, C=1e5)

```

```
logreg = logreg.fit(X_train_word_average, train['tags'])
y_pred = logreg.predict(X_test_word_average)
print('accuracy %s' % accuracy_score(y_pred, test.tags))
print(classification_report(test.tags,
y_pred,target_names=my_tags))
```

```
accuracy 0.6379166666666667
           precision    recall  f1-score   support

     java         0.63      0.59      0.61         613
      html         0.73      0.76      0.75         620
   asp.net         0.65      0.67      0.66         587
        c#         0.53      0.52      0.52         586
ruby-on-rails      0.70      0.77      0.73         599
      jquery      0.44      0.39      0.41         589
      mysql         0.65      0.61      0.63         594
        php         0.73      0.80      0.76         610
        ios         0.60      0.61      0.61         617
  javascript      0.56      0.52      0.54         587
      python      0.55      0.50      0.52         611
         c         0.61      0.61      0.61         594
         css         0.65      0.65      0.65         619
    android         0.60      0.57      0.59         574
     iphone         0.70      0.71      0.71         584
         sql         0.42      0.42      0.42         578
objective-c         0.68      0.71      0.70         591
        c++         0.76      0.78      0.77         608
   angularjs      0.82      0.83      0.82         638
        .net         0.66      0.71      0.68         601

 avg / total         0.63      0.64      0.64       12000
```

Figure 10

It was disappointing, worst we have seen so far.

Doc2vec and Logistic Regression

The same idea of word2vec can be extended to documents where instead of learning feature representations for words, we learn it for sentences or documents. To get a general idea of a word2vec, think of it as a mathematical average of the word vector representations of all the words in the document. Doc2Vec extends the idea of word2vec, however words can only capture so much, there are times when we need relationships between documents and not just words.

The way to train doc2vec model for our Stack Overflow questions and tags data is very similar with when we train Multi-Class Text Classification with Doc2vec and Logistic Regression.

First, we label the sentences. Gensim's Doc2Vec implementation requires each document/paragraph to have a label associated with it. and we do this by using the `TaggedDocument` method. The format will be "TRAIN_i" or "TEST_i" where "i" is a dummy index of the post.

```

1  from tqdm import tqdm
2  tqdm.pandas(desc="progress-bar")
3  from gensim.models import Doc2Vec
4  from sklearn import utils
5  import gensim
6  from gensim.models.doc2vec import TaggedDocument
7  import re
8
9  def label_sentences(corpus, label_type):
10     """
11     Gensim's Doc2Vec implementation requires each document to be a list of sentences.
12     We do this by using the TaggedDocument method. The method takes a list of sentences and a dummy index of the post.
13     a dummy index of the post.
14     """
15     labeled = []
16     for i, v in enumerate(corpus):
17         label_sentences

```

According to [Gensim doc2vec tutorial](#), its doc2vec class was trained on the entire data, and we will do the same. Let's have a look what the tagged document looks like:

```
all_data[:2]
```

```

[TaggedDocument(words=['fulltext', 'search', 'php', 'pdo', 'returning', 'result', 'searched', 'lot', 'matter', 'find', 'wrong',
'setup', 'trying', 'fulltext', 'search', 'using', 'pdo', 'php', 'get', 'results', 'error', 'messages', 'table', 'contains', 'cu
stomer', 'details', 'id', 'int', 'id', 'auto_increment', 'name', 'varchar', '150', 'lastname', 'varchar', '150', 'company', 'va
rchar', '250', 'address', 'varchar', '150', 'postcode', 'int', '5', 'city', 'varchar', '150', 'email', 'varchar', '250', 'phon
e', 'varchar', '20', 'orgnr', 'varchar', '15', 'timestamp', 'timestamp', 'current_timestamp', 'run', 'sqlquery', 'alter', 'tabl
e', 'system_customer', 'add', 'fulltext', 'name', 'lastname', 'except', 'columns', 'id', 'postcode', 'timestamp', 'signs', 'tro
uble', 'far', 'idea', 'problem', 'lies', 'db', 'configuration', 'php', 'code', 'goes', 'php', 'sth', 'dbprepare', 'select', 'n
ame', 'lastname', 'company', 'address', 'city', 'phone', 'email', 'orgnr', 'db_pre', 'customer', 'match', 'name', 'lastname', 'c
ompany', 'address', 'city', 'phone', 'email', 'orgnr', 'search', 'boolean', 'mode', 'bind', 'placeholders', 'sthbindparam', 'sea
rch', 'data', 'sthexecute', 'rows', 'sthfetchall', 'testing', 'print_r', 'dberrorinfo', 'empty', 'rows', 'echo', 'else', 'ech
o', 'foreach', 'rows', 'row', 'echo', 'tr', 'datahref', 'new_orderphp', 'cid', 'row', 'id', 'echo', 'td', 'row', 'name', 'td',
'echo', 'td', 'row', 'lastname', 'td', 'echo', 'td', 'row', 'company', 'td', 'echo', 'td', 'row', 'phone', 'td', 'echo', 'td',
'row', 'email', 'td', 'echo', 'td', 'date', 'ymd', 'strtotime', 'row', 'timestamp', 'td', 'echo', 'tr', 'echo', 'tried', 'chang
e', 'parameter', 'searchquery', 'string', 'like', 'testcompany', 'somename', 'boolean', 'mode', 'also', 'read', 'word', 'foun
d', '50', 'rows', 'counts', 'common', 'word', 'pretty', 'sure', 'case', 'uses', 'specific', 'words', 'table', 'uses', 'mysam',
'engine', 'get', 'results', 'error', 'messages', 'please', 'help', 'point', 'wrong', 'thank', 'tags=[Train_0]]],
TaggedDocument(words=['select', 'everything', '1', 'table', 'x', 'rows', 'another', 'im', 'making', 'join', 'query', 'like',
'select', 'clothes', 'c', 'join', 'style', 'cstyleid', 'ssyleid', 'clothesid', '19', 'dont', 'want', 'select', 'everything',
'style', 'want', 'select', 'everything', 'clothes', '20', 'rows', 'select', '1', 'row', '10', 'style', 'easyst', 'way', 'witho
ut', 'select', 'every', 'row', 'clothes', '20', 'things', 'select', 'like', 'select', 'cid', 'cdescription', 'cname', 'csize',
'cbrand', 'sname', 'clothes', 'c', 'join', 'style', 'cstyleid', 'sstyleid', 'clothesid', '19', 'would', 'fastest', 'way', 'pos
sibility'], tags=[Train_1]]])

```

Figure 11

When training the doc2vec, we will vary the following parameters:

- `dm=0` , distributed bag of words (DBOW) is used.
- `vector_size=300` , 300 vector dimensional feature vectors.
- `negative=5` , specifies how many “noise words” should be drawn.

- `min_count=1` , ignores all words with total frequency lower than this.
- `alpha=0.065` , the initial learning rate.

We initialize the model and train for 30 epochs.

```

1  model_dbow = Doc2Vec(dm=0, vector_size=300, negative=5,
2  model_dbow.build_vocab([x for x in tqdm(all_data)])
3
4  for epoch in range(30):
5      model_dbow.train(utils.shuffle([x for x in tqdm(all
6      model_dbow.alpha -= 0.002
train_doc2vec

```

Next, we get vectors from trained doc2vec model.

```

1  def get_vectors(model, corpus_size, vectors_size, vect
2      """
3      Get vectors from trained doc2vec model
4      :param doc2vec_model: Trained Doc2Vec model
5      :param corpus_size: Size of the data
6      :param vectors_size: Size of the embedding vectors
7      :param vectors_type: Training or Testing vectors
8      :return: list of vectors
9      """
10     vectors = np.zeros((corpus_size, vectors_size))
11     for i in range(0, corpus_size):
12         prefix = vectors_type + ' ' + str(i)
get_vectors

```

Finally, we get a logistic regression model trained by the doc2vec features.

```

logreg = LogisticRegression(n_jobs=1, C=1e5)
logreg.fit(train_vectors_dbow, y_train)
logreg = logreg.fit(train_vectors_dbow, y_train)
y_pred = logreg.predict(test_vectors_dbow)
print('accuracy %s' % accuracy_score(y_pred, y_test))
print(classification_report(y_test,
y_pred, target_names=my_tags))

```

accuracy 0.8045					
	precision	recall	f1-score	support	
java	0.73	0.68	0.70	589	
html	0.89	0.91	0.90	661	
asp.net	0.93	0.94	0.94	606	
c#	0.80	0.80	0.80	613	
ruby-on-rails	0.83	0.90	0.86	601	
jquery	0.72	0.71	0.72	585	
mysql	0.87	0.81	0.84	621	
php	0.81	0.84	0.82	587	
ios	0.68	0.67	0.67	560	
javascript	0.69	0.63	0.66	611	
python	0.63	0.65	0.64	593	
c	0.81	0.83	0.82	581	
css	0.81	0.77	0.79	608	
android	0.84	0.85	0.84	593	
iphone	0.84	0.82	0.83	592	
sql	0.68	0.65	0.66	597	
objective-c	0.84	0.86	0.85	604	
c++	0.90	0.95	0.92	610	
angularjs	0.93	0.96	0.95	595	
.net	0.81	0.84	0.82	593	
avg / total	0.80	0.80	0.80	12000	

Figure 12

We achieve an accuracy score of 80% which is 1% higher than SVM.

BOW with Keras

Finally, we are going to do a text classification with Keras which is a Python Deep Learning library.

The following code were largely taken from a Google workshop. The process is like this:

- Separate the data into training and test sets.
- Use `tokenizer` methods to count the unique words in our vocabulary and assign each of those words to indices.
- Calling `fit_on_texts()` automatically creates a word index lookup of our vocabulary.
- We limit our vocabulary to the top words by passing a `num_words` param to the tokenizer.
- With our tokenizer, we can now use the `texts_to_matrix` method to create the training data that we'll pass our model.
- We feed a one-hot vector to our model.
- After we transform our features and labels in a format Keras can read, we are ready to build our text classification model.
- When we build our model, all we need to do is tell Keras the shape of our input data, output data, and the type of each layer. keras will look after the rest.

- When training the model, we'll call the `fit()` method, pass it our training data and labels, batch size and epochs.

```
1  import itertools
2  import os
3
4  %matplotlib inline
5  import matplotlib.pyplot as plt
6  import numpy as np
7  import pandas as pd
8  import tensorflow as tf
9
10 from sklearn.preprocessing import LabelBinarizer, LabelEncoder
11 from sklearn.metrics import confusion_matrix
12
13 from tensorflow import keras
14 from keras.models import Sequential
15 from keras.layers import Dense, Activation, Dropout
16 from keras.preprocessing import text, sequence
17 from keras import utils
18
19 train_size = int(len(df) * .7)
20 train_posts = df['post'][:train_size]
21 train_tags = df['tags'][:train_size]
22
23 test_posts = df['post'][train_size:]
24 test_tags = df['tags'][train_size:]
25
26 max_words = 1000
27 tokenize = text.Tokenizer(num_words=max_words, char_level=False)
28 tokenize.fit_on_texts(train_posts) # only fit on train
29
30 x_train = tokenize.texts_to_matrix(train_posts)
31 x_test = tokenize.texts_to_matrix(test_posts)
32
33 encoder = LabelEncoder()
34 encoder.fit(train_tags)
35 y_train = encoder.transform(train_tags)
36 y_test = encoder.transform(test_tags)
37
38 num_classes = np.max(y_train) + 1
39 y_train = utils.to_categorical(y_train, num_classes)
```

keras_training

```

Train on 25200 samples, validate on 2800 samples
Epoch 1/2
25200/25200 [=====] - 11s 442us/step - loss: 1.0261 - acc: 0.7180 - val_loss: 0.6658 - val_acc: 0.7975
Epoch 2/2
25200/25200 [=====] - 11s 434us/step - loss: 0.5675 - acc: 0.8190 - val_loss: 0.6625 - val_acc: 0.7868

```

Figure 13

The accuracy is:

```

score = model.evaluate(x_test, y_test,
                       batch_size=batch_size, verbose=1)
print('Test accuracy:', score[1])

```

```

12000/12000 [=====] - 1s 76us/step
Test accuracy: 0.7955833333333333

```

Figure 14

So, which model is the best for this particular data set? I will leave it to you to decide.

Jupyter notebook can be found on Github. Have a productive day!

References:

[https://github.com/RaRe-Technologies/movie-plots-by-genre/blob/master/ipynb with output/Document%20classification%20with%20word%20embeddings%20tutorial%20-%20with%20output.ipynb](https://github.com/RaRe-Technologies/movie-plots-by-genre/blob/master/ipynb%20with%20word%20embeddings%20tutorial%20-%20with%20output.ipynb)

<https://github.com/tensorflow/workshops/blob/master/extras/keras-bag-of-words/keras-bow-model.ipynb>

<https://datascience.stackexchange.com/questions/20076/word2vec-vs-sentence2vec-vs-doc2vec>

